

- Towards a Neural Statistician
- Harrison Edwards Amos Storkey

Idea

Here instead of dealing with a single dataset, we are presented with multiple datasets, and our goal is to obtain a summary for each of these dataset.

Here I use my own notation. Datasets are represented as X_i 's, and their latents are Z_i 's. This is the common assumption used in general generative models. This work introduces a new random variable c called **context**, that accounts for the variations in different datasets. In particular, the model is defined as

$$p(X) = \int p(c) \left[\int p(X|Z)p(Z|c)dZ \right] dc$$

The crucial thing is that there is only a single c for a single dataset.

Basic Model

For the following I will use the paper's notation. They just replace X with D . So their model is

$$p(D) = \int p(c) \left[\prod_{x \in D} \int p(x|z; \theta)p(z|c; \theta)dz \right] dc$$

Inference is two folds. Two inference models are defined:

- $q(c|D; \phi)$ infers the context given the dataset
- $q(z|x, c; \phi)$ infers z given a single x and the context c .

These naturally translates into this lower bound:

$$\mathcal{L}_D = \mathbb{E}_{q(c|D; \phi)} \left[\sum_{x \in d} \mathbb{E}_{q(z|c, x; \phi)} [\log p(x|z; \theta)] - D_{KL}(q(z|c, x; \phi) \| p(z|c; \theta)) \right]$$

Full Model




Figure 1: *Left*: basic hierarchical model, where the plate encodes the fact that the context variable c is shared across each item in a given dataset. *Center*: full neural statistician model with three latent layers z_1, z_2, z_3 . Each collection of incoming edges to a node is implemented as a neural network, the input of which is the concatenation of the edges' sources, the output of which is a parameterization of a distribution over the random variable represented by that node. *Right*: The statistic network, which combines the data via an exchangeable statistic layer.

This just introduces a more powerful generative model, as shown in the middle of figure 1. k latents are introduced, and

$$p(D) = \int p(c) \prod_{x \in D} \int p(x|c, z_{1:L}; \theta) p(z_L|c; \theta) \prod_{i=1}^{L-1} p(z_i|z_{i+1}, c; \theta) dz_{1:L} dc$$

Inference:

$$q(c, z_{1:L}|D; \phi) = q(c|D; \phi) \prod_{x \in D} q(z_L|x, c; \phi) \prod_{i=1}^{L-1} q(z_i|z_{i+1}, x, c; \phi)$$

Just trivial.

Statistical Network

We call the network that models $q(c|D; \phi)$ as a statistical network. In this work, this involves

- An instance encoder E that encodes individual datapoint x_i to $e_i = E(x_i)$
- An exchangeable instance polling layer that collpases the matrix (e_1, \dots, e_k) to a single pre-statistic vector v .
- A final post polling network that takes v to a parametrization of a diagonal Gaussian.

Experiments

The most important thing is to note that, we must have a **large number** of datasets. They are preferred to be small.

- Simple 1-D distributions: each dataset is a single type of distribution of various mean and variance. The context describes the type, mean and variance of the distribution
- Spatial MNIST: each single image is represented as a set of points, and thus a single dataset.
- Omniglot. The important thing is that is composes of several classes. Three examples of fea-shot learning:
 - Training on a subset of datasets, and do conditional generation (dataset reconstruction) on unseen datasets in OMNIGLOT
 - Do generation on MNIST

- Few-shot classification. The idea is just to use c as a class embedding.
- Youtube faces: same as above.