

- Variational Continual Learning
- Cuong V. Nguyen, E. Turner

The Problem

We are given a sequence of datasets D_1, \dots, D_T . This dataset can even be for different tasks. Our goal is, upon receiving a new dataset D_t , update our model from m_{t-1} , to m_t , so that it adapts to the new dataset, but **still works for all previous datasets**. This is continual learning.

The Idea

Assume that datasets and datapoints are independent (but they can have different distributions, although they may share parameters). Then, $p(\theta|D_{1:T})$ can be decomposed as

$$p(\theta|D_{1:T}) \propto p(D_{1:T}|\theta)p(\theta) = p(D_{1:T-1}|\theta)p(D_T|\theta)p(\theta) = p(\theta|D_{1:T-1})p(D_T|\theta)$$

Here, $p(\theta|D_{1:T})$ will be our model that summarizes $D_{1:T}$. When given D_{T+1} , we compute $p(\theta|D_{1:T+1})$ by multiplying $p(\theta|D_{1:T})$ with $p(D_T|\theta)$, and normalizing.

However, $p(\theta|D_{1:T})$ is intractable. So we will use an $q_t(\theta)$ to approximate this. If we define the operator *proj* as

$$proj(p^*(\theta)) = \arg \min_{q \in Q} KL(q(\theta) \| p^*(\theta))$$

Then we can compute $q_t(\theta)$ recursively as

$$q_t(\theta) = proj(q_{t-1}(\theta)p(D_t|\theta))$$

Episodic Memory Enhancement

$q_t(\theta)$ is an approximation. Plus, q_t is computed recursively. This means error can easily accumulate. To mitigate this problem, we extend VCL to include a small representative set of data C_t for $D_{1:t}$, which we will call the "coreset". **We hope the information from this coreset will be passed to the final approximation as precise as possible.** We will decompose

$$p(\theta|D_{1:T}) = p(\theta|D_{1:T} \setminus C_t)p(C_t|\theta) = \tilde{q}_t(\theta)p(C_t|\theta)$$

In this way, only $p(\theta|D_{1:T} \setminus C_t)$ will be computed recursively. In particular, it will be computed as follows:

$$p(\theta|D_{1:T} \setminus C_t) = p(\theta|D_{1:t-1} \setminus C_{t-1})p(C_{t-1} \setminus C_t|\theta)p(D_t \setminus C_t|\theta) \approx \tilde{q}_{t-1}(\theta)p(D_t \cup C_{t-1} \setminus C_t|\theta)$$

You can just draw a Venn graph to see why this is true.

Algorithm 1 Coreset VCL

Input: Prior $p(\theta)$.

Output: Variational and predictive distributions at each step $\{q_t(\theta), p(y^*|\mathbf{x}^*, \mathcal{D}_{1:t})\}_{t=1}^T$.

Initialize the coreset and variational approximation: $C_0 \leftarrow \emptyset, \tilde{q}_0 \leftarrow p$.

for $t = 1 \dots T$ **do**

Observe the next dataset \mathcal{D}_t .

$C_t \leftarrow$ update the coreset using C_{t-1} and \mathcal{D}_t .

Update the variational distribution for non-coreset data points:

$$\tilde{q}_t(\theta) \leftarrow \arg \min_{q \in \mathcal{Q}} \text{KL}(q(\theta) \parallel \frac{1}{Z} \tilde{q}_{t-1}(\theta) p(\mathcal{D}_t \cup C_{t-1} \setminus C_t | \theta)). \quad (2)$$

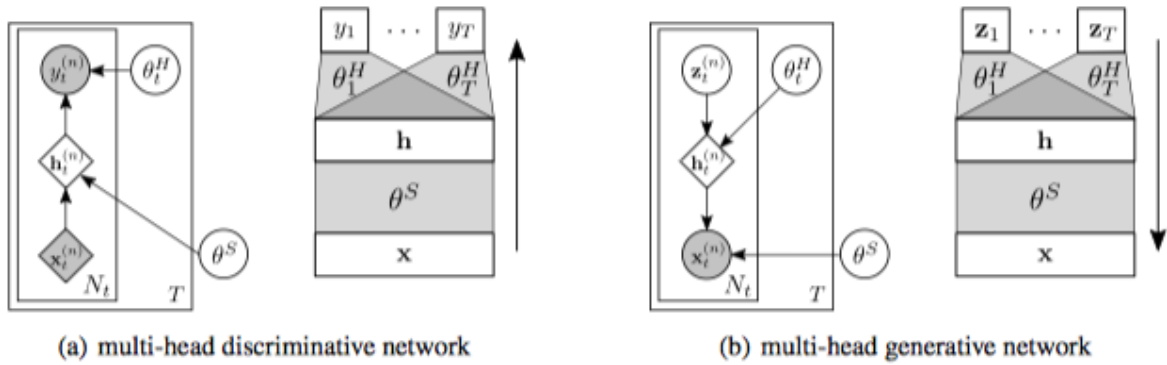
Compute the final variational distribution (only used for prediction, and not propagation):

$$q_t(\theta) \leftarrow \arg \min_{q \in \mathcal{Q}} \text{KL}(q(\theta) \parallel \frac{1}{Z} \tilde{q}_t(\theta) p(C_t | \theta)). \quad (3)$$

Perform prediction at test input \mathbf{x}^* : $p(y^*|\mathbf{x}^*, \mathcal{D}_{1:t}) = \int q_t(\theta) p(y^*|\theta, \mathbf{x}^*) d\theta$.

end for

VCL in Discriminative Models



There are two cases:

1. All datapoints come from the same distribution
2. Datapoints in different datasets come from different distributions, **and even the output can differ**.

For the first case, we can just use a single-head network. For the second, there will be multiple heads.

Note, here we will define different $p(y|x, \theta)$ for different tasks. θ will include a shared part θ^S , but will also include task specific part θ_t^H . This weights are only used in their $p(y|x, \theta)$, which means that they will be updated only once.

VCL in Generative Models

We assume a single $p(z)$. So $p(x|z, \theta)$ must differ for different tasks. Other things are just trivial.

Experiments

Dicriminative

- Permuted-MINST: trivial. Single head-network
- Split-MNIST. 0/1, 2/3, 4/5. We must use multi-head network

Generative

- MNIST: 0, 1, 2, 3, ...

- Characters.